

Project AURA: System Architecture & Presentation Guide

An Ambient, Edge-Native Personal AI Companion with Streaming Windowing, Autonomous Denoising, Multi-Task Extraction, and Medallion Long-Term Memory.

Author: Makeathon
Engineering Team

Domain: Data
Engineering & Edge AI

Target OS: Android 14+
(Local NPU/CPU)

Security: Zero Cloud
Audio Egress

The 30-Second Elevator Pitch for Judges

*"Judges, cloud-based assistants like Siri or ChatGPT have two fatal flaws for continuous ambient use: **Privacy leakage** (streaming your room's microphone to external servers 24/7) and **unsustainable cloud cost/latency**."*

***Aura is an ambient, 100% on-device AI companion.** It continuously processes audio streams using a data-engineering-inspired **Dual-Trigger Window (60s time-cap or 30s silence)**, runs autonomous acoustic error correction using an embedded **Gemma 2B INT4 LLM**, extracts actionable commitments, and organizes knowledge into a local **Medallion Data Architecture (Bronze → Silver → Gold)**. Zero raw audio ever leaves the phone. Zero cloud API bills. Complete offline privacy."*

1. The Problem & Engineering Innovation

Wearable ambient devices (like OMI / formerly Friend, Limitless Pendant, and Rewind) have validated huge consumer demand for an AI that remembers everything you hear and say. However, virtually all current market solutions stream raw audio to cloud APIs (Whisper + GPT-4o).

✗ Cloud-Centric Ambient Assistants


- **Severe Privacy Risk:** Enterprise meetings, NDAs, passwords, and private conversations stream to 3rd-party servers.
- **High Operating Cost:** Continuous audio streaming costs ~\$0.06/min. An 8-hour workday costs **\$800+/month per user**.
- **Network Fragility:** Dies in elevators, underground subways, airplanes, or low-connectivity zones.

✓ Aura: Edge-Native Paradigm

- **Air-Gapped Privacy:** 100% on-device inference via Google MediaPipe GenAI & Gemma 2B. Audio never touches the network.
- **Zero Marginal Cost:** \$0.00 cloud inference bill. Runs on consumer phone hardware (CPU NEON / GPU / NPU).
- **Local Medallion Storage:** Ephemeral raw speech buffer auto-prunes in 24 hours, distilling into compact long-term memories.

2. High-Level Design (HLD)

Aura is structured as an **Event-Driven Streaming Architecture** that treats ambient speech as a continuous, unbounded data stream.



2.1 The Data Engineer's Lens: Concept Translation Table

Presenting this table will immediately resonate with technical judges, showcasing architectural maturity beyond a typical hackathon prototype:

Data Engineering Concept	Enterprise Big Data Primitive	Project AURA Edge Equivalent	Benefit / Rationale
Continuous Streaming Ingestion	Kafka Topic / Event Hub	<code>SpeechManager</code> continuous microphone listener	Zero disk buffering for audio; immediate real-time transcription.
Windowing & Watermarking	Apache Flink Tumbling/Session Windows	Dual-Trigger: 60s hard cap OR 30s silence timeout	Bounds memory consumption; triggers inference only on coherent thoughts.
Bronze Data Layer	Raw Delta Lake / Iceberg table	<code>ConversationSegment</code> Room DB table	Audit trail of raw spoken words; 24-hour auto-prune retention TTL.
Silver Cleaned Layer	Denoised / Deduplicated view	Gemma 2B Contextual Speech Denoising pass	Fixes phonetic ASR errors (e.g., "date and friendship" → "data pipeline").
Gold Aggregated Layer	Curated Business Feature Store	<code>Memory</code> & <code>Task</code> database entities	Micro-compacted permanent knowledge (~120 tokens/thought) for instant recall.
Backpressure / Concurrency	Kafka Consumer Group Rate Limiter	Coroutine <code>Mutex.withLock</code> Gatekeeper	Prevents multi-threaded native C++ segmentation faults (SIGSEGV).

3. Low-Level Design (LLD): Component Architecture

AuraForegroundService.kt — Unkillable Ambient Service

1

Registers a persistent Android `FOREGROUND_SERVICE_MICROPHONE` with a low-priority notification channel. Bypasses Android 14 aggressive Doze/App Standby battery killing, allowing Aura to listen in the background with the screen turned off.

SpeechManager.kt — Zero-Chime Continuous Audio Listener

2

Encapsulates Android `SpeechRecognizer` with `EXTRA_PREFER_OFFLINE = true`. Solves the infamous Android mic start-beep by dynamically muting `STREAM_MUSIC` and `STREAM_SYSTEM` during listener resets, ensuring completely silent 24/7 background operation.

AuraViewModel.kt — Stateful Stream Processor & Window Engine

3

Coordinates asynchronous flows. Manages the 60s hard-cap timer (`maxBlockDurationJob`) and 30s silence watchdog (`silenceWatchdogJob`). Aggregates conversational chunks into a single UI card in real-time, preventing the "split dialogue card" bug.

GemmaManager.kt — Local LLM Inference Engine & Task/Memory Core

4

Direct JNI interface to Google MediaPipe's `tasks-genai` runtime loading `gemma-2b-it-cpu-int4.bin`. Houses the thread-safe `inferenceMutex` and implements prompt pipelines for Denoising, Multi-Task Extraction, and Memory Compaction.

AuraDatabase.kt & DAOs — Edge Storage & Pruning Engine

5

SQLite Room implementation with optimized indices. Implements `deleteOldSummarized(beforeTimestamp)` to maintain a constant database footprint (< 20 MB) regardless of how many months the user runs the application.

4. Core Algorithms & Logic Flow

4.1 The Dual-Trigger Windowing Algorithm (60s Cap vs 30s Silence)

Ambient conversations contain natural pauses, interruptions, and long soliloquies. A naive timer either splits a sentence in half or waits indefinitely. Aura's dual-trigger mechanism

provides mathematical bounds on latency and memory:

```
ON SpeechSegmentCaptured(text):
    now = CurrentTimestamp()

    // 1. Check if the existing conversation block hit the 60-second time cap
    if (now - blockStartTimestamp ≥ 60,000 ms):
        FinalizeCurrentBlock(reason = "Max 60s Reached")

    // 2. If starting a new block, set start timestamp and schedule 60s timer
    if (blockStartTimestamp == 0):
        blockStartTimestamp = now
        ScheduleTimer(MAX_DURATION = 60,000 ms) → FinalizeCurrentBlock("60s Cap")

    // 3. Append speech to live dialogue card (immediate visual feedback)
    currentBlockBuilder.append(text)
    UpdateLiveTranscriptCard(currentBlockBuilder.toString())

    // 4. Reset the 30-Second Inactivity Watchdog
    CancelTimer(silenceWatchdogTimer)
    silenceWatchdogTimer = ScheduleTimer(SILENCE_TIMEOUT = 30,000 ms) →
        FinalizeCurrentBlock("30s Silence")
```

4.2 Autonomous ASR Denoising & Homophone Correction

When a phone sits across a desk, ambient acoustic distortion causes phonetic errors (e.g. *"date and friendship"* instead of *"data pipeline"*, *"iceberk"* instead of *"iceberg"*). Upon block finalization, Gemma 2B runs an in-place semantic restoration pass:

```
user
You are an intelligent ambient speech corrector. Spoken English was recorded by a phone mic

Tasks:
1. Fix obvious speech recognition mistakes into sensible, coherent natural English based on
2. Keep the exact meaning and original words where they make sense. Do NOT add new facts or
3. If the transcription already makes good sense, keep it as is.
4. Output ONLY the cleaned transcript with no preamble or explanation.

Raw speech: "$rawText"
Cleaned text:
model
```

4.3 Multi-Task Extraction (Hybrid Pattern + LLM Engine)

To guarantee zero dropped commitments even during heavy model loads, Aura executes a hybrid extraction strategy:

- **Phase 1 (Deterministic Look-Ahead):** A regex engine scans for triggers (`remind me to` , `action item` , `don't forget to` , `todo`) and isolates action verbs from meeting context.
- **Phase 2 (Temporal Extraction):** Dedicated regex extracts deadlines (`\b(\d{1,2}(\?:\d{2})?\s*(?:am|pm)|tomorrow|monday)\b`).
- **Phase 3 (Gemma Formalization):** When Gemma is ready, it transforms multiple commitments into structured imperative task objects, deduplicating via signature hashes (`title_dueTime`).

5. Technical Specifications & Benchmark Data

 Model & Quantization Model: Google Gemma 2B Instruct Quant: INT4 (4-bit integer weights) Storage Size: 1.52 GB on disk Context Window: 2,048 tokens	 Runtime & Compute Runtime: MediaPipe GenAI 0.10.22 Execution: ARM CPU NEON / Qualcomm Hexagon NPU RAM Usage: ~1.8 GB during inference Inference Speed: ~1.2s for 150-word block	 Battery & Thermal Idle State: Mic DSP (0.5% battery/hr) Active Window: Wake-on-window (2s CPU burst) Thermal Profile: Negligible rise (< 1.5°C over 1 hr meeting)
---	--	--

6. Judge Q&A Defense Strategy (Winning the Q&A)

Q1: "Why not use a Cloud LLM like GPT-4o mini? It's smarter and supports longer contexts."

Answer: *"In an ambient listening application, cloud LLMs are a non-starter for two reasons: First is **Privacy & Compliance**. Employees cannot wear a device that streams confidential enterprise meetings, client health records, or proprietary code to third-party cloud servers. Second is **Unit Economics**. Streaming 8 hours of audio per day to Whisper and cloud LLMs costs ~\$25/day per user (\$750/month). Aura runs 100% on-device for \$0.00 marginal cost."*

Q2: "How does the system prevent the database from growing out of control if it listens all day?"

Answer: *"We implement a **Medallion Data Architecture with automated compaction and TTL (Time-To-Live) pruning**. Raw Bronze conversation segments are ephemeral and auto-pruned after 24 hours. Only distilled Gold memories (~120 tokens per summary) and structured Tasks are permanently stored in SQLite. Even after months of daily meetings, the entire database consumes less than 20 MB."*

Q3: "What happens if two tasks are spoken quickly in the same sentence?"

Answer: *"Our hybrid task engine splits the text stream across consecutive trigger boundaries (e.g., 'remind me to X' and 'also action item to Y'). Each action item is extracted, matched against temporal deadline expressions (e.g., 'by 3 PM' or 'before Friday'), deduplicated via signature hashing, and inserted as separate records into the Room database."*

Q4: "Doesn't continuous microphone listening drain the user's phone battery?"

Answer: *"We use a **Tiered Compute Architecture**. Continuous audio capture is offloaded to the phone's dedicated audio DSP running Android's native speech recognizer at negligible power draw (~0.5%/hr). The high-compute Gemma 2B model is kept completely asleep, waking up only for a 2-second burst when a 60-second window completes or 30 seconds of silence is reached. The CPU remains idle 95% of the time."*

